

Clustering Metagenome Sequences Using Canopies

Mohammad Arifur Rahman, Nathan LaPierre, Huzefa Rangwala and Daniel Barbara

Department of Computer Science

George Mason University

Fairfax, VA, United States

Email: mrahma23@gmu.edu, nlapier2@gmu.edu, rangwala@cs.gmu.edu and dbarbara@gmu.edu

Abstract—The advancement in genome technologies has allowed for the collective sequencing of co-existing microbial communities (Metagenomics), omnipresent across various environments like the soil, ocean and human body. Metagenomics has spurred the development of several bioinformatics approaches for analyzing the diversity, abundance, function and role of the different organisms within these communities.

We present an approximate but fast clustering algorithm for analyzing large-scale metagenomic data. Our approach achieves efficiency by first partitioning the large number of sequence reads into groups (called canopies) using a fast locality sensitive hashing based distance function. This initial approximate assignment of sequence reads to canopies is refined by using state-of-the-art sequence clustering algorithms. This two-phase approach allows for the use of our developed algorithms as a pre-processing phase for computationally expensive clustering algorithms. Using the clusters as surrogates for Operational Taxonomic Units (OTUs) we can estimate the biodiversity within a community sample.

The Canopy-based clustering algorithm is evaluated on synthetic and real world 16S and whole metagenome benchmarks. We demonstrate the ability of our proposed approach to determine meaningful OTU assignments and observe significant speedup with regards to run time when combined with three different state-of-the-art sequence clustering algorithms.

Keywords—Clustering, Canopy, Metagenome, 16S, Biodiversity

I. INTRODUCTION

A large portion of the earth’s biomass comprises trillions of tiny microorganisms with varying biodiversity and plays a crucial role in ensuring the stability of environment. Communities of interacting microbes exists in several ecosystems varying from the soil, ocean and human body [1]; and though, most of them are beneficial to the host some are known to cause unwanted conditions and can be linked to cause of diseases in the host [2].

Metagenomics is the sequencing of the collective DNA of microbial organisms coexisting as communities. Analyzing metagenomes has provided an unprecedented opportunity to understand the diversity, role and function of these organisms within clinical and ecological environments [3][4].

State-of-the-art genome sequencing technologies do not deliver the complete genome of an organism (millions in length), but large number of short contiguous subsequences called *reads* of length 75 to 500 (based on technology) in

random order. Sequencing communities of genomes exacerbates the problem since reads from different microbes are mixed together and upfront the composition of a community in terms of abundance and identity of bacteria is unknown [5]. Further, sequencing technologies produce datasets that range from Gigabytes (GB) to Terabytes (TB) with their own idiosyncrasies. Alternatively, targeted metagenome sequencing that involves sequencing of *marker genes* has been popular for the characterization of these communities. 16S and 18S sequences are repetitive marker genes within microbial genomes but have variations in their sequences that allow them to be separated into different taxonomic groups [6].

Over the years, several unsupervised clustering algorithms have been developed and used for the analysis of large-scale targeted and whole metagenome sequences (reviewed in Section II). The unsupervised grouping of similar sequences from the 16S/18S genes is referred as *binning* and allows for the estimation of biodiversity within a sample. Binning leads to assignment of similar sequences within groups called *Operational Taxonomic Units (OTUs)* [7]. Microbial OTUs are generally ecologically consistent across the hosts regardless of OTU clustering approaches [8].

In this paper we develop an approximate but efficient clustering algorithm based on Canopy Clustering [9] which can be used as a pre-clustering step with any state-of-the-art sequence clustering algorithms. Specifically, our approach identifies canopies with a greedy procedure and a fast sequence distance metric based on locality sensitive hashing [10]. Sequences within the canopies are considered as an initial partition (or grouping) of data which can be further refined by applying expensive and accurate clustering algorithms. Our developed approach treats each canopies independently allowing for easy parallelization.

We present experimental results on real and synthetic metagenome sequence benchmarks and show that use of Canopy Clustering (CC) in combination with UCLUST [11], SUMACLUSt [13] and SWARM [12] leads to improved run time efficiency by maximum of 12.19, 21.09 and 18.61 times respectively with accurate clustering results. We also show the performance of our developed approach is scalable for large datasets with increased number of processors. The

source code is freely available on Github¹ for public use.

II. LITERATURE REVIEW

Over the years several sequence clustering methods have been developed and used widely for metagenome sequences. A comprehensive survey by Kopylova et. al. [14] benchmarks various sequence clustering algorithms including UCLUST [11], SWARM [12], SUMACLUSt [13] and MOTHUR [15]

CD-HIT [16] is general purpose sequence clustering algorithm that follows an incremental, greedy approach. CD-HIT uses pairwise sequence alignment to find similar sequences. UCLUST [11] is similar to CD-HIT but achieves a significant speedup over CD-HIT by using seeds (fixed length gapless subsequences) for performing pairwise sequence comparisons. MC-LSH [17] utilizes an efficient locality sensitive based hashing function to approximate the pairwise sequence similarity. MC-MinH [18] uses min-wise [19] hashing along with the greedy clustering to group 16S and whole metagenome sequences. Mash [20], uses MinHash locality sensitive hashing to reduce large sequences to a representative sketch and estimate pairwise distances. Other methods for clustering sequence reads include TOSS [21], AbundanceBin [22] and CompostBin [23]. All unique kmers are first grouped in TOSS and then clusters are merged based on kmer repetitions. In AbundanceBin, reads are modeled as a mixture of Poisson distributions. Expectation Maximization (EM) algorithm is used to infer model parameters for the final clustering. Principal component analysis is used within CompostBin to project the data into a lower dimensional space, followed with a graph partitioning approach. MOTHUR [15] uses a pairwise distance matrix as input and performs hierarchical clustering. An expensive sequence alignment is used between all pairs of input sequences for computing distances.

SWARM [12] uses exhaustive single-linkage clustering based on optimal sequence alignment. Sequences that are less than a certain distance from any other other sequence in the cluster are clustered together. SWARM attempts to reduce the impact of clustering parameters on the resulting OTUs by avoiding arbitrary global clustering thresholds and input sequence ordering dependence. At first SWARM builds an initial set of OTUs is constructed by iteratively agglomerating similar amplicons. Then amplicon abundance values are used to reveal OTUs internal structures and to break them into sub-OTUs.

SUMACLUSt [13] follows similar approach as UCLUST. Based on greedy strategy SUMACLUSt incrementally constructs clusters by comparing an abundance-ordered list of input sequences against the representative set of already-chosen sequences. Initially this list is empty. UCLUST and recently developed, SWARM and SUMACLUSt are

considered to be the state-of-the-art metagenome sequence clustering methods by a benchmarking study [14].

In this study we developed an approximate but highly efficient sequence clustering algorithm using Canopies [9]. The developed approach can be used as a pre-processing step for other state-of-art sequence clustering algorithms to improve efficiency and also retains correctness in clustering.

III. METHODS

A. Overview

Figure 1 shows an overview of our proposed canopy clustering approach for metagenome sequences. Sequence reads are represented with *kmers* defined as contiguous subsequence of length k . We select a read randomly as the canopy centroid. Then we compute the distances of other reads to this canopy centroid. Two thresholds are used in Canopy clustering algorithm. If the a distance is less than a *tight* threshold (T_2) we assign the read to the canopy and that read will not be considered for another canopy. Otherwise, if a distance is less than a *soft* threshold (T_1) we assign the read to the canopy but that read will be available for another canopy. This process continues until all sequence reads is assigned to at least one canopy. For fast canopy assignment we use a random projection based technique. Specifically, we use Locality Sensitive Hash (LSH) to compute pairwise distances. More accurate clustering methods are used to sub-cluster each canopy in parallel. Sub-clustering method considers only the members within the canopy which reduces pairwise distance calculations.

B. Canopy Clustering

Canopy Clustering [9] is an efficient approximate clustering algorithm often used as pre-processing step for other accurate and expensive clustering methods like the K-means algorithm or the Hierarchical clustering algorithm. It is intended to speed up clustering operations on large data sets, where standard clustering algorithms may be impractical in terms of run time and memory consumption due to the large number of pairwise distance calculations. For a dataset with N instances, worst case calculations without canopy clustering is N^2 . After canopy clustering, if the number of canopies is k then worst case calculations with canopy clustering is $\sum_{i=1}^k (c_i)^2$ where c_i is the number of instances within i th canopy.

Canopy clustering uses two distance thresholds, (i) *soft* threshold T_1 and (ii) *tight* threshold T_2 . If data point p_1 is within the soft distance threshold T_1 with centroid p_2 then p_1 will reside in same canopy as p_2 but p_1 may belong to other canopies assuming that it has only met soft threshold and best match is yet to be found. Thus one data point may belong to multiple canopies in Canopy clustering. On the other hand if data point p_1 is within the tight distance threshold T_2 with centroid p_2 then canopy clustering assigns p_1 to the same canopy as p_2 and stops assigning p_1 to

¹<https://github.com/mrahma23/LSH-Canopy>

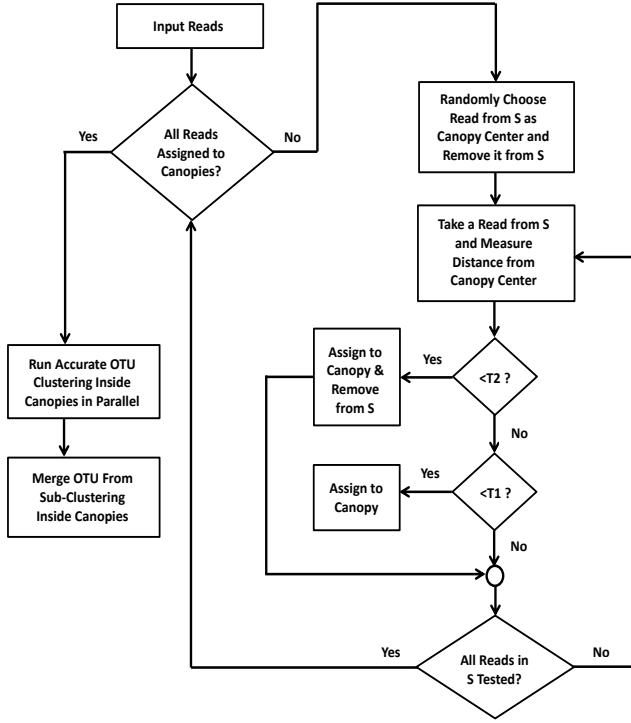


Figure 1. Workflow of Canopy Clustering for Large Scale Metagenome Data

any other canopy assuming that tight threshold has been met and best canopy assignment for p_1 has been found. Canopy centroids are selected randomly until all data points is assigned to at least one canopy.

C. Locality Sensitive Hashing

Canopies are intended to reduce pairwise distance calculations. Canopy clustering itself should be efficient which requires a fast and approximate distance measure for Canopy assignments. Locality Sensitive Hashing (LSH) [10] is an algorithm for solving the approximate or exact near neighbors search. Suppose we have a space S of points with a distance measure $d(x, y)$. A family H of hash functions is said to be (d_1, d_2, p_1, p_2) -sensitive for any x and y in S : if $d(x, y) < d_1$, then the probability over all $h \in H$, that $h(x) = h(y)$ is at least p_1 and if $d(x, y) > d_2$, then the probability over all $h \in H$, that $h(x) = h(y)$ is at most p_2 . LSH provides a fast approximate distance measure while reducing data dimensionality which makes LSH appropriate for Canopy clustering in 16S and whole metagenomic data.

We constructed the LSH family with bit sampling [24]. The normalized kmers frequency based feature vectors of sequence reads were first projected into d dimensional vectors in $\{0, 1\}^d$ space. Given an input vector v and a random hyperplane defined by r , we let $h(v) = \{0, 1\}$ based on $sgn(v \cdot r) = \pm 1$ that indicates on which side of the

hyperplane v lies. Each possible choice of r defines a single function. Let H be the set of all such functions. For any $h_i \in H$ and for any two data x, y the probability that x and y agree on i^{th} positions of their respective d -length binary vector is

$$P[h_i(x) = h_i(y)] = 1 - \frac{distance(h_i(x), h_i(y))}{d} \quad (1)$$

Where $distance$ is the hamming distance and d is the number of bits. Hence $H = \{h_1, h_2, \dots, h_d\}$ is a $(d_1, d_2, 1-d_1/d, 1-d_2/d)$ sensitive LSH family. Both the random projection and hamming distance calculation is cheap and fast which makes this scheme appropriate for our purpose which is fast partitioning of large scale genome sequences so that more expensive clustering can be done in parallel.

D. Sub-Clustering Inside Canopies

Canopy clustering makes initial approximate partitions of the dataset and reduces pairwise distance computations. Each of these partitions can be further sub clustered in parallel with more accurate and expensive clustering methods with lower computational cost since only the points inside canopies are taken into consideration. We have used three recent and popular sequence clustering methods as more accurate and expensive sub-clustering measure inside canopies in this study. UCLUST [11], SUMACLUSt [13] and SWARM [12] were used for sub-clustering canopies.

E. Merging results from Canopies

Each cluster (OTU) is represented by the Longest Common Subsequence of all member sequences. The final step of our proposed framework is to merge OTU representations generated by canopies. According to Canopy cluster algorithm a single data point may belong to multiple canopies as long as the soft threshold is met. As a result similar OTU representations may appear from multiple canopies. To eliminate redundancy we run UCLUST on the OTU representations.

In the following Sections UCLUST, SUMACLUSt and SWARM with Canopy clustering approach are represented as CC_{UCLUST} , $CC_{SUMACLUSt}$ and CC_{SWARM} respectively where the term CC stands for Canopy Clustering.

IV. EXPERIMENTAL EVALUATION

A. Dataset Description

To evaluate the performance of our proposed approach we use previously published synthetic and real world 16S, 18S and metagenome sequence benchmarks. Key statistics and relevant information regarding these datasets are presented in Table I.

Synthetic Datasets:

Table II
PERFORMANCE COMPARISON [*F*-SCORE AND PEARSON CORRELATION COEFFICIENT (ρ)]

Methods	Comparison Metric	Datasets							
		Synthetic			Real World				
		Bokulich ₂	Bokulich ₃	Bokulich ₆	Body Sites	Canadian Soil	Global Soil	Liver Cirrhosis Metagenome (Sampled 3M)	Liver Cirrhosis Metagenome (Sampled 30M)
UCLUST	<i>F</i> -Measure	0.39	0.40	0.51	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>
<i>CC</i> _{UCLUST}	<i>F</i> -Measure	0.39	0.41	0.52	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>
	(ρ) with Respect to UCLUST	0.9831	0.9753	0.9831	0.9682	0.8419	0.9824	0.9216	0.9072
SUMACLUSt	<i>F</i> -Measure	0.40	0.41	0.51	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>
<i>CC</i> _{SUMACLUSt}	<i>F</i> -Measure	0.41	0.42	0.51	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>
	(ρ) with Respect to SUMACLUSt	0.9709	0.9813	0.9538	0.9518	0.7643	0.8714	0.9281	0.8614
SWARM	<i>F</i> -Measure	0.46	0.48	0.55	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>
<i>CC</i> _{SWARM}	<i>F</i> -Measure	0.46	0.49	0.56	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>	<i>N/A</i>
	(ρ) with Respect to SWARM	0.9817	0.97861	0.9251	0.9648	0.7581	0.9143	0.9263	0.8533

Table shows values of *F*-measures and Pearson Correlation Coefficient (ρ -value) of UCLUST, SUMACLUSt, SWARM and their respective versions with Canopy clustering. *F*-measures is only available for synthetic datasets but not for real world datasets since no ground truths like known taxonomic profiles are available for them. ρ -value was calculated based on the taxonomy profiles at Genus level generated from clustered OTUs provided by a method and it's corresponding Canopy counterpart. Higher *F*-measures reflect better clustering by adhering to ground truth. Higher ρ -values reflect stronger correlation between taxonomic profiles. Higher *F*-scores and ρ values are represented with bold and italic letters.

Table III
BIODIVERSITY COMPARISON [FAITHS PHYLOGENETIC DIVERSITY METRIC (PD), SHANNON AND SIMPSON]

Methods	Comparison Metric	Datasets							
		Synthetic			Real World				
		Bokulich ₂	Bokulich ₃	Bokulich ₆	Body Sites	Canadian Soil	Global Soil	Liver Cirrhosis Metagenome (Sampled 3M)	
UCLUST	PD Range	[171.95 – 221.85]	[186.90 – 212.84]	[104.51 – 104.51]	[1.46 – 46.79]	[0.30 – 1352.73]	[2.98 – 3.29]	[3.14 – 51.47]	
	Shannon Range	[2.52 – 3.51]	[2.43 – 3.54]	[5.87 – 5.87]	[0.29 – 7.67]	[2.32 – 7.86]	[1.84 – 8.30]	[0.27 – 5.77]	
	Simpson Range	[0.55 – 0.75]	[0.55 – 0.76]	[0.96 – 0.96]	[0.049 – 0.98]	[0.80 – 0.99]	[0.0 – 0.98]	[0.02 – 0.97]	
<i>CC</i> _{UCLUST}	PD Range	[164.79 – 217.72]	[169.41 – 198.36]	[109.33 – 109.33]	[2.37 – 47.13]	[0.52 – 1419.31]	[3.02 – 3.81]	[4.61 – 53.29]	
	Shannon Range	[2.61 – 3.83]	[2.92 – 3.91]	[6.41 – 6.41]	[0.93 – 7.13]	[3.26 – 7.61]	[3.72 – 7.38]	[0.13 – 6.17]	
	Simpson Range	[0.64 – 0.87]	[0.56 – 0.93]	[0.96 – 0.96]	[0.081 – 0.99]	[0.84 – 0.99]	[0.21 – 0.99]	[0.09 – 0.98]	
SUMACLUSt	PD Range	[106.00 – 162.78]	[142.85 – 174.19]	[89.22 – 89.22]	[0.93 – 39.47]	[0.59 – 1279.29]	[2.98 – 3.29]	[0.28 – 49.61]	
	Shannon Range	[2.00 – 3.01]	[2.19 – 3.28]	[5.48 – 5.48]	[0.16 – 7.43]	[2.32 – 7.32]	[1.00 – 7.89]	[1.37 – 7.83]	
	Simpson Range	[0.52 – 0.73]	[0.54 – 0.75]	[0.95 – 0.95]	[0.027 – 0.98]	[0.80 – 0.99]	[0.40 – 0.98]	[0.19 – 0.91]	
<i>CC</i> _{SUMACLUSt}	PD Range	[114.96 – 171.57]	[147.85 – 187.91]	[93.81 – 93.81]	[0.86 – 41.63]	[0.81 – 1292.34]	[1.37 – 4.89]	[0.18 – 51.35]	
	Shannon Range	[2.51 – 3.94]	[2.96 – 4.11]	[5.94 – 5.94]	[1.21 – 7.13]	[3.12 – 7.79]	[2.17 – 7.25]	[1.91 – 7.39]	
	Simpson Range	[0.68 – 0.79]	[0.51 – 0.74]	[0.96 – 0.96]	[0.06 – 0.99]	[0.88 – 0.99]	[0.23 – 0.98]	[0.27 – 0.88]	
SWARM	PD Range	[18.37 – 24.73]	[17.36 – 19.81]	[30.84 – 30.84]	[1.44 – 28.66]	[0.54 – 706.57]	[5.79 – 6.18]	[2.06 – 39.71]	
	Shannon Range	[2.98 – 3.91]	[2.01 – 3.04]	[5.03 – 5.03]	[0.28 – 7.63]	[1.0 – 7.79]	[1.66 – 7.81]	[1.47 – 6.91]	
	Simpson Range	[0.70 – 0.82]	[0.53 – 0.74]	[0.95 – 0.95]	[0.05 – 0.98]	[0.50 – 0.99]	[0.00 – 0.99]	[0.18 – 0.89]	
<i>CC</i> _{SWARM}	PD Range	[19.18 – 26.87]	[18.43 – 22.61]	[31.48 – 31.48]	[2.34 – 29.97]	[1.37 – 748.71]	[2.34 – 8.46]	[3.18 – 40.73]	
	Shannon Range	[1.66 – 4.87]	[1.19 – 4.13]	[6.03 – 6.03]	[0.89 – 7.13]	[2.81 – 8.06]	[2.81 – 7.87]	[2.03 – 6.88]	
	Simpson Range	[0.66 – 0.88]	[0.41 – 0.81]	[0.91 – 0.91]	[0.11 – 0.99]	[0.74 – 0.99]	[0.14 – 0.99]	[0.07 – 0.86]	

Table shows ranges of values for Faiths Phylogeny Diversity (PD), Shannon and Simpson coefficient over all samples in a dataset in the format [*minimum* – *maximum*]. Most of these datasets contain multiple samples and Alpha diversity metrics like PD, Shannon and Simpson values are generated for each of these samples separately. Biodiversity metric values changes significantly over samples e.g. diversity from hair samples and teeth cavity are supposedly different. So instead of mean values this Table represents [*minimum* – *maximum*] ranges of values a sample can take. Similar ranges reflect similar diversity.

1) *Bokulich₂*: This dataset was prepared using the Illumina TruSeq v2 paired-end library preparation kit. It is a simulated 16S rRNA gene microbial community dataset. This dataset contains 19 taxonomic Families, 19 Genera, 22 Species and 22 Strains in total. This dataset can also be found in the QIIME database (identifier 1685).

2) *Bokulich₃*: Similar to *Bokulich₂* except that it was prepared with the TruSeq v1 paired-end library kit at Illumina Cambridge and is also available in the QIIME database (identifier 1686).

3) *Bokulich₆*: This 16S rRNA dataset was sequenced at Washington University School of Medicine and contains

evenly distributed microbial communities. This dataset contains 13 taxonomic Families, 23 Genera, 44 Species and 48 Strains in total.

All these datasets from Bokulich et al. [26] are available at QIIME database² under their respective ID's. Since, these are simulated datasets the taxonomic profile of microbial organisms within them are known.

Real World Datasets:

4) *Canadian Soil*: The Canadian Soil dataset³ contains genomic data of soil spanning from Arctic Tundra to Agri-

²http://qiime.org/home_static/dataFiles.html

³<http://www.cm2bl.org/>

Table IV
RUNTIME COMPARISON (IN MINUTES)

Datasets			Methods								
Type	Title	# of Reads	UCLUST	CC _{UCLUST}	Speed Up	SUMACLUSt	CC _{SUMACLUSt}	Speed Up	SWARM	CC _{SWARM}	Speed Up
Synthetic	Bokulich ₂	6,938,836	12.71	3.08	4.13x	114.53	13.89	8.24x	128.12	17.24	7.43x
	Bokulich ₃	3,594,237	10.43	3.61	2.89x	27.73	5.11	5.43x	18.37	3.39	5.41x
	Bokulich ₆	250,903	4.47	2.09	2.14x	5.61	2.07	2.71x	4.17	1.29	3.21x
	Body Sites	886,630	9.46	3.03	3.12x	18.42	6.02	3.06x	16.37	5.83	2.81x
Real World	Canadian Soil	2,966,053	13.65	4.17	3.27x	363.96	56.61	6.43x	117.53	20.84	5.64x
	Global Soil	9,252,764	108.21	18.75	5.77x	510.92	45.37	11.26x	289.51	34.84	8.31x
	Liver Cirrhosis Metagenome	3,000,000	14.57	4.03	3.61x	46.62	9.02	5.17x	41.37	8.75	4.73x
	Liver Cirrhosis Metagenome	30,000,000	22.41h	1.84h	12.19x	46.62h	2.21h	21.09x	37.43h	2.01h	18.61x

Table I
DATASET STATISTICS

Datasets	Type	# of Reads	# of Samples	Platform
Bokulich ₂ [26]	M	6,938,836	4	H
Bokulich ₃ [26]	M	3,594,237	4	H
Bokulich ₆ [26]	M	250,903	1	H
Canadian Soil [27]	R	2,966,053	13	H
Body Sites [28]	R	886,630	602	G
Global Soil [29]	R	9,252,764	57	H
Liver Cirrhosis [30]	R	6,117,828,130	232	H

Table shows information about dataset used in this study. M, R, H and G represent Mock, Real World, HiSeq and GS-FLX, respectively.

cultural soil suitable for different agricultural products.

5) *Body Sites*: This dataset contains composition of bacterial communities from up to 27 different body sites in healthy adults. A collection of 602 samples acquired from different body sites of human subjects are provided with meta-data.

6) *Global Soil*: The global soil data was taken from Ramirez et al. [29] which is a study of the below-ground diversity in New York City's Central Park.

7) *Liver Cirrhosis*: The Liver Cirrhosis dataset was taken from the study [30] which is a whole gut microbiome wide association study of stool samples from 98 liver cirrhosis patients and 83 healthy controls to characterize the fecal microbial communities and their functional composition. In total, 860 Gb of high-quality sequence data was generated in this study. Because of the high volume of sequence reads in this dataset, we used random sampling for feasible clustering performance evaluations. Three million and thirty million sequence reads were randomly chosen from original dataset regardless of sample labels (disease or control) for our study.

B. Evaluation Metrics

We evaluate the performance of our developed clustering approach using the following commonly used metrics that are used for the assessment of (i) outputs from clustering algorithms, (ii) biodiversity within metagenome samples and (iii) computational run time.

1) *Faiths phylogenetic diversity metric (PD)*: Faiths phylogenetic diversity [31] combines all the branch lengths of phylogenetic tree as a measure of diversity. So, if a new OTU is found and it is closely related to another OTU in the sample, it will contribute to a small increase to the PD score.

However, if a new OTU from different lineage is found then it will contribute to a large increase in the PD score.

2) *Shannon Entropy*: Shannon-Wiener diversity index is defined as:

$$H = - \sum_{i=1}^s (p_i \log_2 p_i) \quad (2)$$

where s is the number of OTUs and p_i is the proportion of the community represented by OTU i . The Shannon index increases as both the richness and evenness of the community increase. The fact that the index incorporates both components of biodiversity can be seen as both a strength and a weakness. It is a strength because it provides a simple summary, but it is a weakness because it makes it difficult to compare communities that differ greatly in richness.

3) *Simpson's Index*: Simpsons index is defined as $1 - \text{dominance}$ or

$$1 - \sum p_i^2 \quad (3)$$

where where p_i is the proportion of the community represented by OTU i . Simpsons index is based on the probability that any two individuals drawn at random from an infinitely large community belong to the same species. It measures *evenness* of the community from 0 to 1. Higher value of this index implies higher similarity and relatively lower diversity of microorganisms within a sample.

4) *F-measure*: In case of synthetic datasets, expected taxonomic composition is known (*ground truth*). False-positive (FP) refers to the number of taxonomy that was found in observed but not expected, false-negative (FN) refers to the number of taxonomy that exists in expected but not observed, and true-positive (TP) refers to the number of taxonomy exists in both observed and expected. The following definitions were used:

$$\text{precision} = \frac{TP}{(TP + FP)} \quad (4)$$

$$\text{recall} = \frac{TP}{(TP + FN)} \quad (5)$$

$$FScore = \frac{2 \times \text{precision} \times \text{recall}}{(\text{precision} + \text{recall})} \quad (6)$$

5) *Pearson Coefficient Correlation (ρ -value)*: After getting Operational Taxonomic Units (OTU) from a clustering method we create a taxonomic profile at the Genus level. Pearson's correlation coefficient was computed to measure the relatedness of taxonomic assignment between a pair of tools. Values range between -1 and 1, with -1 indicating a negative correlation, 0 indicating no correlation, and 1 indicating a positive correlation or strong relationship.

C. Parameter Settings

For kmers, the value of parameter k was set to 4. The parameters for bit length of LSH (d), canopy's soft ($T1$) and hard ($T2$) thresholds were set by using random sampling and validation. For validation purpose, only ten percent of data was randomly sampled from the dataset of interest. First the parameter for LSH bit length d was estimated. In order to know correctness of our estimation of d , we performed Canopy clustering with corresponding sub clustering inside canopies and compared the results corresponding to the bit lengths. For synthetic dataset $F1$ -scores and for real world datasets Pearson Coefficient Correlation (ρ) with results from corresponding expensive naive sequence clustering. We started with bit length $d=1$ and continued to increase value of d as long as we got better results for sampled data. Initial value of soft ($T1$) and tight ($T2$) threshold was set to 0.6 and 0.4 respectively where maximum distance is 1.0. These are relaxed initial values for both $T1$ and $T2$, meaning that they allow for higher repetitions of sequence reads in multiple canopies due to higher value of $T1$ and easy final canopy assignment due to higher value of $T2$ which results in rough and fast approximate clusters prior to expensive sub-clustering step. Any bit length d that brings comparatively better results for these relaxed $T1$ and $T2$ on sampled data is expected to bring better results for more strict $T1$ and $T2$ on whole datasets.

Once d is selected we estimated values for $T1$ and $T2$ on the same sampled data. First, hard threshold ($T2$) was estimated with soft threshold ($T1$) fixed at 0.6. $T2$ determines which sequence reads will be retained for next iteration of canopy clustering. Same comparison metrics from the estimation of parameter d were used. We decreased the value of $T2$ from 0.4 (*relaxed*) to 0.1 (*strict*) with a step size of 0.01. The $T2$ corresponding to best result was selected. Then $T2$ was fixed at the best estimation and $T1$ was varied from 0.6 to best estimation for $T2$ with step size of 0.01. The value for $T1$ corresponding to best result was selected. Table V shows the estimated parameter values for our study.

D. Hardware and Software for Experiments

We performed all the experiments on computers with Intel 5th generation Core i7 2.70GHz 64bit processor with 8 core CPUs and 12GB memory. Similarity identity of sub-clustering methods were set to default 97%. For implemen-

Table V
PARAMETER SETTINGS

Datasets	Total # of Reads	# of Reads in Sampled Data	Parameters		
			d	$T1$	$T2$
Bokulich ₂	6,938,836	693,884	47	0.48	0.36
Bokulich ₃	3,594,237	359,428	31	0.43	0.34
Bokulich ₆	250,903	25,090	17	0.37	0.21
Canadian Soil	2,966,053	296,605	29	0.46	0.37
Body Sites	886,630	88,663	22	0.39	0.22
Global Soil	9,252,764	925,276	59	0.49	0.38
Liver Cirrhosis	3,000,000	300,000	48	0.42	0.35
Liver Cirrhosis	30,000,000	3,000,000	67	0.51	0.39

tation we used Python 2.7.12 and QIIME [32] version 1.9.0. Taxonomy for reported OTUs was assigned using the RDP Classifier [33] against the 97% representative databases for Greengenes [34] and Silva [35] for methods used in this study. We used PyNast⁴ open source sequence aligner for aligning clustered output.

V. RESULTS DISCUSSION

A. Clustering Performance

Table II shows the performance of UCLUST, SUMACLUSt and SWARM with their corresponding versions with CC clustering. Table II compares F-scores and Pearson Correlation Coefficient (ρ). F-scores are only available for synthetic benchmarks since taxonomic profile for them is known as ground truth. Correlation values were generated based on taxonomic profiles at Genus level generated from outputs from clustering methods. From Table II we can see that F-scores obtained from a clustering method and its corresponding Canopy clustering version are very similar and in some cases better. We see a higher F-score for Bokulich₂ benchmark from CC_{SUMACLUSt}. CC_{UCLUST} and CC_{SWARM} provided same F-scores as their respective naive versions for Bokulich₂. For Bokulich₃ benchmark we observed higher F-scores for all Canopy clustering methods. Finally for Bokulich₆ benchmark F-scores of CC_{UCLUST} and CC_{SWARM} were improved comparing to respective naive versions.

For all benchmarks we observe strong correlations between taxonomic profiles at genus level. The highest correlation was observed for Bokulich₆ benchmark between UCLUST and CC_{UCLUST} with 0.9831. From observed F-scores and correlation values we can say that running a sequence clustering method with our proposed LSH based Canopy cluster framework can bring similar or better results.

B. Biodiversity Estimation

Clustering metagenome sequences outputs OTUs that represent biodiversity contained in the samples from which data is collected. We compare the biodiversity represented by OTUs from naive methods and their respective Canopy clustering variations used in this study. Table III shows

⁴<http://biocore.github.io/pynast/>




Figure 2. Effect of number of processors on Runtime of CC_{UCLUST} , $CC_{SUMACLUST}$ and CC_{SWARM} for Bokulich₂ and GLocal Soil datasets, two of the largest datasets used in this study.

Faiths phylogenetic diversity metric (PD), Shannon and Simpson index after clustering with different methods. These metrics are some of the popular Alpha Diversity metrics meaning that they measure species diversity in sites or habitats at a local scale. These metric values are generated per sample basis. Table III shows ranges of metric values in $[minimum - maximum]$ format provided by methods with and without Canopy clustering. Any sample of a dataset will take value from this range. These ranges may not be same since OTUs vary over clustering methods. But they should be similar for a method and it's Canopy counterpart. We can see from Table III that Canopy clustering based methods produce similar ranges of values as their naive counterparts. No significant changes in diversity metric values were observed which indicates that Canopy based approaches reproduce similar biodiversity while reducing run time.

C. Runtime Comparison

Table IV shows the run time in minutes of UCLUST, SUMACLUST, SWARM and their respective versions with our proposed Canopy clustering pipeline. We see from Table IV that CC_{UCLUST} outperforms UCLUST for all datasets. The highest speed up for UCLUST was observed for the metagenome data containing 30 million randomly sampled sequence reads where CC_{UCLUST} was 12.19 times faster than UCLUST. $CC_{SUMACLUST}$ outperforms SUMACLUST in all cases. The highest speed-up for $CC_{SUMACLUST}$ was observed to be 21.09x for the large metagenome dataset. The highest speedups for CC_{SWARM} was also observed for the large metagenome dataset with 30 million reads. From these observations we conclude that Canopy clustering not only reduces the run time of an expensive clustering method but also scales well for larger benchmark datasets comparing to the smaller ones.

D. Effect of Varying Number of Processors

Figure 2a-2b shows the runtime of CC_{UCLUST} , $CC_{SUMACLUST}$ and CC_{SWARM} on two largest benchmarks titled Bokulich₂ and Global Soil. We observe that




Figure 3. 3a and 3b show effect of varying $T1$ and $T2$ on F -scores for the largest synthetic dataset Bokulich₂

increasing the number of processors reduces the total runtime. Steeper curves show reduction in runtime with increasing number of processors. Significant reductions in run time were observed $CC_{SUMACLUST}$ and $CC_{SUMACLUST}$ compared to CC_{UCLUST} .

E. Sensitivity Analysis (Varying $T1$ and $T2$)

Figure 3a-3b shows effect of varying $T1$ ($T2$ fixed at 0.34) and $T2$ ($T1$ fixed at 0.45) on F -scores for the largest synthetic dataset Bokulich₂. Reducing $T1$ leads to comparatively *strict* soft-threshold which will reduce repetitions of instances in multiple canopies. For a fixed $T2$ this implies that lower $T1$ will yield better canopy assignment. From Figure 3a we can say that when $T1$'s range is in 0.4 to 0.6 our proposed approach provides better F -Scores. On the other hand, increasing $T2$ leads to comparatively *relaxed* tight-threshold for canopies. As a result instances will be prematurely assigned to canopies without waiting for best match. From Figure 3b we can say that when $T2$'s range is in 0.25 to 0.35 our proposed approach achieves better F -Scores. Lower $T2$ may lead to higher runtime since instances will continue to reappear until the $T2$ is met.

VI. CONCLUSION AND FUTURE WORK

We developed a greedy approximate clustering process for any accurate and relatively expensive clustering on large scale metagenome datasets. Our approach takes advantage of the multi-core CPU systems by partitioning the large dataset with a fast and cheap pairwise distance measure and then deploying comparatively expensive clustering in parallel which considers only data points that are inside a partition. Our proposed approach scales well with large datasets and provide significant reduction in computation time. We demonstrate that our approach provides similar outcome in terms of biodiversity metrics, ground truth and taxonomic correlation with corresponding expensive clustering methods.

REFERENCES

- [1] Turnbaugh et al., "The human microbiome project: exploring the microbial part of ourselves in a changing world," *Nature*, vol. 449, no. 7164, p. 804, 2007.

- [2] P. J. Turnbaugh, R. E. Ley, M. Hamady, C. Fraser-Liggett, R. Knight, and J. I. Gordon, "The human microbiome project: exploring the microbial part of ourselves in a changing world," *Nature*, vol. 449, no. 7164, p. 804, 2007.
- [3] Qin et al., "A human gut microbial gene catalogue established by metagenomic sequencing," *nature*, vol. 464, no. 7285, pp. 59–65, 2010.
- [4] M. Pop and S. L. Salzberg, "Bioinformatics challenges of new sequencing technology," *Trends in Genetics*, vol. 24, no. 3, pp. 142–149, 2008.
- [5] H. Teeling and F. O. Glöckner, "Current opportunities and challenges in microbial metagenome analysis: a bioinformatic perspective," *Briefings in bioinformatics*, p. bbs039, 2012.
- [6] Chakravorty et al., "A detailed analysis of 16s ribosomal rna gene segments for the diagnosis of pathogenic bacteria," *Journal of microbiological methods*, vol. 69, no. 2, pp. 330–339, 2007.
- [7] Blaxter et al., "Defining operational taxonomic units using dna barcode data," *Philosophical Transactions of the Royal Society B: Biological Sciences*, vol. 360, no. 1462, pp. 1935–1943, 2005.
- [8] T. S. Schmidt, J. F. M. Rodrigues, and C. Von Mering, "Ecological consistency of ssu rna-based operational taxonomic units at a global scale," *PLoS Comput Biol*, vol. 10, no. 4, p. e1003594, 2014.
- [9] McCallum, Andrew and Nigam, Kamal and Ungar, Lyle H., "Efficient clustering of high-dimensional data sets with application to reference matching," in *Proceedings of the Sixth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, ser. KDD '00. New York, NY, USA: ACM, 2000, pp. 169–178.
- [10] A. Gionis, P. Indyk, and R. Motwani, "Similarity search in high dimensions via hashing," in *Proceedings of the 25th International Conference on Very Large Data Bases*, ser. VLDB '99. San Francisco, CA, USA: Morgan Kaufmann Publishers Inc., 1999, pp. 518–529.
- [11] Edgar, Robert C., "Search and clustering orders of magnitude faster than blast," *Bioinformatics*, vol. 26, no. 19, pp. 2460–2461, 2010.
- [12] Mahé et al., "Swarm v2: highly-scalable and high-resolution amplicon clustering," *PeerJ*, vol. 3, p. e1420, 2015.
- [13] C. Mercier, F. Boyer, A. Bonin, and E. Coissac, "Sumatra and sumacust: fast and exact comparison and clustering of sequences [submitted for publication]," Available at <https://git.metabarcoding.org/obitools/sumatra/wikis/home>, 2014.
- [14] Kopylova et al., "Open-source sequence clustering methods improve the state of the art," *mSystems*, vol. 1, no. 1, 2016.
- [15] Schloss et al., "Introducing mothur: Open-source, platform-independent, community-supported software for describing and comparing microbial communities," *Applied and Environmental Microbiology*, vol. 75, no. 23, pp. 7537–7541, 2009.
- [16] Li, Weizhong and Godzik, Adam, "Cd-hit: a fast program for clustering and comparing large sets of protein or nucleotide sequences," *Bioinformatics*, vol. 22, no. 13, pp. 1658–1659, 2006.
- [17] Z. Rasheed, H. Rangwala, and D. Barbara, "Efficient clustering of metagenomic sequences using locality sensitive hashing," in *SDM. SIAM*, 2012, pp. 1023–1034.
- [18] Z. Rasheed and H. Rangwala, *MC-MinH: Metagenome Clustering using Minwise based Hashing*, ch. 74, pp. 677–685.
- [19] A. Z. Broder, M. Charikar, A. M. Frieze, and M. Mitzenmacher, "Min-wise independent permutations," *Journal of Computer and System Sciences*, vol. 60, no. 3, pp. 630–659, 2000.
- [20] Ondov et al., "Mash: fast genome and metagenome distance estimation using minhash," *bioRxiv*, p. 029827, 2016.
- [21] O. Tanaseichuk, J. Borneman, and T. Jiang, "Separating metagenomic short reads into genomes via clustering," *Algorithms for Molecular Biology*, vol. 7, no. 1, p. 1, 2012.
- [22] Wu, Yu-Wei and Ye, Yuzhen, "A novel abundance-based algorithm for binning metagenomic sequences using l-tuples," *Journal of Computational Biology*, vol. 18, no. 3, pp. 523–534, 2011.
- [23] S. Chatterji, I. Yamazaki, Z. Bai, and J. A. Eisen, "Compost-bin: A dna composition-based algorithm for binning environmental shotgun reads," in *Annual International Conference on Research in Computational Molecular Biology*. Springer, 2008, pp. 17–28.
- [24] P. Indyk and R. Motwani, "Approximate nearest neighbors: towards removing the curse of dimensionality," pp. 604–613, 1998.
- [25] A. Z. Broder, "On the resemblance and containment of documents," pp. 21–29, 1997.
- [26] Bokulich et al., "Quality-filtering vastly improves diversity estimates from illumina amplicon sequencing," *Nat Methods*, vol. 10, pp. 57–59, 2013.
- [27] Neufeld et al., "Open resource metagenomics: a model for sharing metagenomic libraries," *Standards in Genomic Sciences*, vol. 5, pp. 203–210, 2011.
- [28] Costello et al., "Bacterial community variation in human body habitats across space and time," *Science*, vol. 326, no. 5960, pp. 1694–1697, 2009.
- [29] Ramirez et al., "Biogeographic patterns in below-ground diversity in new york city's central park are similar to those observed globally," *Proceedings of the Royal Society of London B: Biological Sciences*, vol. 281, no. 1795, 2014.
- [30] N. et al., "Alterations of the human gut microbiome in liver cirrhosis," *Nature*, vol. 513, no. 7516, pp. 59–64, 2014.
- [31] Faith, Daniel P, "Conservation evaluation and phylogenetic diversity," *Biological conservation*, vol. 61, no. 1, pp. 1–10, 1992.

- [32] Caporaso et al., “Qiime allows analysis of high-throughput community sequencing data,” *Nat Meth*, vol. 7, no. 5, pp. 335–336, 2010.
- [33] Q. Wang, G. M. Garrity, J. M. Tiedje, and J. R. Cole, “Naive bayesian classifier for rapid assignment of rna sequences into the new bacterial taxonomy,” *Applied and environmental microbiology*, vol. 73, no. 16, pp. 5261–5267, 2007.
- [34] DeSantis et al., “Greengenes, a chimera-checked 16s rna gene database and workbench compatible with arb,” *Applied and environmental microbiology*, vol. 72, no. 7, pp. 5069–5072, 2006.
- [35] Pruesse et al., “Silva: a comprehensive online resource for quality checked and aligned ribosomal rna sequence data compatible with arb,” *Nucleic acids research*, vol. 35, no. 21, pp. 7188–7196, 2007.