

Clustering Metagenomic Sequences Using Canopies

Mohammad Arifur Rahman, Nathan LaPierre, Huzefa Rangwala and Daniel Barbara

Department of Computer Science

George Mason University

Fairfax, VA, United States

Email: mrahma23@gmu.edu, nlapier2@gmu.edu, rangwala@cs.gmu.edu and dbarbara@gmu.edu

Abstract—Metagenomics is the study of genome sequences from samples which are retrieved directly from an environment hosting multiple organisms that co-exist as communities within ecological regions. Metagenomics has created the necessity and challenges of developing computational tools for the quantification of abundance, diversity, role of different species within communities and phenotype inferences. Enormous amount of data is created during metagenome sequencing process. Algorithms have been developed to cluster similar metagenome sequences. We have developed an approach for clustering metagenomic data that uses Canopy Clustering with Locality Sensitive Hashing distance approximation to make the clustering process in metagenomic data faster. Canopy Clustering works as a preprocessing step to reduce pairwise distance calculation and enables efficient parallel processing with subsequent expensive cluster methods while LSH provides fast distance approximation and reduces data dimension. We tested our framework with three popular clustering methods in literature on three synthetic and three real world large scale metagenome datasets and observed that our proposed approach can reduce runtime while providing similar and in some cases better outcomes.

Keywords—Clustering, Canopy, LSH, MST, Metagenome

I. INTRODUCTION

A large portion of Earth’s biomass comprises microorganisms that exhibit remarkable genetic diversity and play important role in Earth’s ecosystem. Microorganisms reside as interacting communities. Existence, interaction and evolution of microorganisms in the hosting environment is crucial to nutrient recycling, nitrogen cycle, precipitation, weather and moreover, environmental stability. But certain microbial abundance and interaction may cause unwanted conditions like disease in hosting sites. Analysis of microbial community can reveal interesting relationships between the microbial community and the host. For example analyzing metagenomic data from human gut microbiome can reveal the role of microbes in human health [2]. Phenotype inferences like which bacteria are responsible for ulcerative colitis and how a patient is most likely to get the abnormal abundance of such bacteria, can be made. Metagenomics has created the opportunity to study genome sequences of microbiomes as interacting communities. It is the study of genetic material recovered directly from environmental samples. These samples can be taken from sea, soil and the human body [1].

State-of-the-art sequencing technologies produce short contiguous subsequences called *reads* from random positions of actual whole genome. In whole metagenome sequencing these reads from different organisms are combined together posing fundamental challenge for further analysis. Similar sequence reads are needed to be clustered prior to further analysis. Metagenome data can be based on targeted 16S and 18S rRNA gene sequences. 16S and 18S sequences are marker genes, which exists in most microbial species but have variations in their sequences that allow them to be separated into different taxonomic groups [5]. These sequences are highly repetitive in metagenome data. Hence clustering is required for binning similar sequences based on similarity among them. These clusters of similar sequences are called Operational Taxonomic Units (OTUs). OTUs are used to classify groups of closely related individuals from similar or different taxonomic levels [6]. Each of these clusters may represent taxonomic units of bacteria depending on the sequence similarity threshold. Microbial OTUs are generally ecologically consistent across the hosts regardless of OTU clustering approaches [7]. For this reason 16S/18S gene sequencing and OTU clustering have been widely used for the analysis of genetic diversity within complex microbial communities.

Metagenome data contains large number of reads, repetitions, ambiguity and missing sequence fragments. These characteristics of the data make OTU clustering a time consuming and challenging task. OTU clustering approaches have been developed for the analysis of microorganisms, some of which are discussed in Section II. Hence improving scalability and efficiency of OTU clustering methods on metagenome data still remain as a challenge and active area of research.

In this study we propose and evaluate a pre-clustering technique based on the Canopy clustering and Locality Sensitive Hashing. Canopy clustering method reduces pairwise distance calculations between sequences by approximately partitioning the data. Only the sequences inside a canopy will be considered for further sub-clustering within that canopy. More accurate and expensive clustering methods can be used for sub-clustering within each partition in parallel. This way our proposed approach can improve the scalability and efficiency of clustering in large scale metagenome data.

II. LITERATURE REVIEW

Over the years several sequence clustering methods have been developed and used widely for metagenomic sequence data. A comprehensive survey by Kopylova et. al. [8] benchmarks various sequence clustering algorithms including UCLUST [9], SWARM [10], SUMACLUSt [11], MOTHUR and SORTMERA [12].

CD-HIT [13] is general purpose sequence clustering algorithm that follows an incremental, greedy approach. CD-HIT uses pairwise sequence alignment to find similar sequences. UCLUST [9] is similar to CD-HIT but achieves a significant speedup over CD-HIT by using seeds (fixed length gapless subsequences) for performing pairwise sequence comparisons. MC-LSH [14] utilizes an efficient locality sensitive based hashing function to approximate the pairwise sequence similarity. MC-MinH [15] uses min-wise [16] hashing along with the greedy clustering to group 16S and whole metagenomic sequences. Mash [17], uses MinHash locality sensitive hashing to reduce large sequences to a representative sketch and rapidly estimate pairwise distances. Other methods for clustering metagenomic sequence reads include TOSS [18], AbundanceBin [19] and CompostBin [20]. All unique k-mers are first clustered in TOSS and then clusters are merged based on k-mer repetitions. In AbundanceBin reads are modeled as mixture of Poisson distributions. Then Expectation Maximization (EM) algorithm is used to infer model parameters for the final clustering. Principal component analysis is used within CompostBin to project the data into a lower dimensional space, followed with a graph partitioning approach. MOTHUR [12] uses a pairwise distance matrix as input and perform hierarchical clustering. An expensive sequence alignment is used between all pairs of input sequences for computing distances.

UCLUST [9] and recently developed, SWARM [10] and SUMACLUSt [11] are considered to be the state-of-the-art metagenomic sequence clustering methods by a benchmarking study [21]. SWARM uses an exhaustive single-linkage clustering approach with a fast pairwise sequence alignment. SUMACLUSt is similar to UCLUST. Abundance ordered list of input sequences are compared against the representative set of already chosen sequences in SUMACLUSt.

In this study we developed an approximate but highly efficient sequence clustering algorithm using canopies [22]. The developed approach can be used as a pre-processing step with other state-of-art sequence clustering algorithms to improve efficiency but retain correctness of the clustering. Specifically, we show the performance of our developed approach in conjunction with UCLUST, SWARM and SUMACLUSt as detailed in Section III-D.

III. METHODS

A. Overview

Figure 1 shows a detailed overview of our proposed Canopy clustering approach for metagenomic data. After




Figure 1. Workflow of Canopy Clustering for Large Scale Metagenome Data

computing the kmers and Locality Sensitive Hash (LSH) codes our proposed method will start canopy clustering. LSH are used as approximate distance for canopy assignments. Canopy clustering will stop when all reads in datasets are assigned to at least one canopy. These canopies represent initial and approximate partitions of the large metagenome data. More accurate clustering methods will be used to accurately cluster each canopy. Sub-clustering inside a canopy will only take the members of that canopy into account which will reduce pairwise distance calculations. This sub-clustering will be performed in parallel. In this study UCLUST, SUMACLUSt, SWARM and Minimum Spanning Tree (MST) will be used as accurate sub-clustering methods for canopies which are briefly discussed in Section III-D.

B. Canopy Clustering

Canopy Clustering [22] is an efficient approximate clustering algorithm often used as preprocessing step for other accurate and expensive clustering methods like the K-means algorithm or the Hierarchical clustering algorithm. It is intended to speed up clustering operations on large data sets, where standard clustering algorithms may be impractical in terms of run time and memory consumptions due to the large number of pairwise distance calculations. For a dataset with N instances, worst case calculations without canopy

clustering is N^2 . After canopy clustering if the number of canopies is k then worst case calculations with canopy clustering is $\sum_{i=1}^k (c_i)^2$ where c_i is the number of instances within i th canopy.

Canopy clustering uses two distance thresholds, (i) *soft* threshold $T1$ and (ii) *tight* threshold $T2$. If data point p_1 is within the soft distance threshold $T1$ with another data point p_2 then p_1 will reside in same canopy as p_2 but p_1 may belong to other canopies assuming that it has only met soft threshold and best match is yet to be found. Thus one data point may belong to multiple canopies in Canopy clustering. On the other hand if data point p_1 is within the tight distance threshold $T2$ with another data point p_2 then canopy clustering assigns p_1 to the same canopy as p_2 and stops assigning p_1 to any other canopy assuming that tight threshold has been met and best canopy assignment for p_1 has been found. In this case p_1 will not be repeated in other canopies. Canopy centroids are selected randomly. Then all other data points are tested for canopy assignment based on $T1$ and $T2$. New canopies are selected until all data points are assigned to canopies. A fast approximation based on Locality Sensitive Hashing (LSH) is used for Canopy membership identification.

C. Locality Sensitive Hashing

Canopies are intended to reduce pairwise distance calculations. Canopy clustering itself should be efficient which requires a fast and approximate distance measure for Canopy assignments. Locality Sensitive Hashing (LSH) [23] is an algorithm for solving the approximate or exact Near Neighbor Search in high dimensional spaces. LSH reduces data dimensionality by randomly projecting data points to low dimensional bit signatures and provides approximate distance measure. These characteristics make LSH appropriate for dimensionality reduction and relatively accurate distance measure for Canopy clustering in metagenome data.

In this study we have chosen a random projection based hashing function that projects a d dimensional data point into a n dimensional bit representation where $n < d$ and reduces dimensionality. Given a random projection v' of size $1 \times d$ and a data vector a of size $1 \times d$, the dot product between them provides a scalar value. Sign of this scalar value represents which side of the random hyperplane does the data point exist. This information can be represented with a single bit. This way a single data point can be represented with limited number of bits. Hamming distance between these binary representations indicates *disagreements* between two data points.

D. Sub-Clustering Inside Canopies

We have used three recent and popular sequence clustering methods as the expensive clustering measure inside canopies in this study. UCLUST [9], SUMACLUSt [11] and SWARM [10] were used for sub-clustering canopies.

1) *UCLUST*: The 32 bit version of UCLUST executable is available for academic usage. But the 64-bit versions which is necessary to handle large datasets, require expensive license. UCLUST follows a greedy process and creates *seeds* of sequences which generate clusters based on percent identity.

2) *SUMACLUSt*: It follows similar approach as UCLUST. Based on greedy strategy SUMACLUSt incrementally constructs clusters by comparing an abundance-ordered list of input sequences against the representative set of already-chosen sequences. Initially this list is empty.

3) *SWARM*: SWARM uses exhaustive single-linkage clustering based on optimal sequence alignment. Sequences that are less than a certain distance from any other sequence in the cluster are clustered together. SWARM attempts to reduce the impact of clustering parameters on the resulting OTUs by avoiding arbitrary global clustering thresholds and input sequence ordering dependence. At first SWARM builds an initial set of OTUs is constructed by iteratively agglomerating similar amplicons. Then amplicon abundance values are used to reveal OTUs internal structures and to break them into sub-OTUs.

4) *Minimum Spanning Tree (MST)*: We also implement Minimum Spanning Tree (MST) based clustering method as an accurate and expensive sub-clustering method within canopies. Minimum Spanning Tree (MST) produces a connected spanning tree of undirected graph. It connects all vertices together such that the total weight of edges remains minimum. This property has been used for clustering [24] [25]. We represent edges with LSH based distances. Then edges with highest distances were removed from MST to generate OTUs.

E. Merging results from Canopies

Each OTU is represented by a single sequence which is the Longest Common Subsequence of all member sequences of that OTU. The final step of our proposed framework is to merge OTU representations generated by canopies. According to Canopy cluster algorithm a single data point may belong to multiple canopies as long as the soft threshold is met. As a result similar OTU representations may appear from multiple canopies. To eliminate redundancy we run UCLUST on the OTU representations. We also use Minimum Spanning Tree and Locality Sensitive Hashing (LSH) distance based clustering method to cluster OTU representations. A brief description of MST based clustering is given in Section III-D. Unlike merging with greedy clustering methods like UCLUST, this approach requires number of edges to remove as prior information.

In the following Sections UCLUST, SUMACLUSt, SWARM and MST with Canopy clustering approach are represented as CC_{UCLUST} , $CC_{SUMACLUSt}$, CC_{SWARM}

and CC_{MST} respectively where the term CC stands for Canopy Clustering.

IV. EXPERIMENTAL EVALUATION

A. Dataset Description

To evaluate the performance of our proposed approach we use previously published synthetic and real world sequence benchmarks. These datasets are publicly available¹. We used three synthetic 16S rRNA gene mock community datasets (Bokulich₂, Bokulich₃, and Bokulich₆) from Bokulich et al. [26] and three real data sets: a 16S rRNA gene soil data set (Canadian Soil) [27], a 16S rRNA gene human data set (Body Sites) [28] and 18S rRNA gene soil data set (Global Soil) [29]. Key statistics and relevant information regarding these datasets are presented in Table I.

Table I
DATASET STATISTICS

Datasets	Type	# of Reads	# of Samples	Read Length	Platform
Bokulich ₂ [26]	M	6,938,836	4	189–251	H
Bokulich ₃ [26]	M	3,594,237	4	114–151	H
Bokulich ₆ [26]	M	250,903	1	114–150	H
Canadian Soil [27]	R	2,966,053	13	76–10	H
Body Sites [28]	R	886,630	602	117–351	F
Global Soil [29]	R	9,252,764	57	119–151	H

Table shows information about dataset used in this study. M, R, H and F represent Mock, Real World, HiSeq and GS-FLX respectively.

Synthetic Datasets:

1) *Bokulich₂*: This dataset was prepared using the Illumina TruSeq v2 paired-end library preparation kit. It is a simulated 16S rRNA gene microbial community data set. This dataset contains 19 taxonomic Families, 19 Genera, 22 Species and 22 Strains in total. This dataset can be found in the QIIME database (identifier 1685).

2) *Bokulich₃*: Similar to Bokulich₂ except that it was prepared with the TruSeq v1 paired-end library kit at Illumina Cambridge and is also available in the QIIME database (identifier 1686).

3) *Bokulich₆*: This 16S rRNA dataset was sequenced at Washington University School of Medicine and contains evenly distributed microbial communities. This dataset contains 13 taxonomic Families, 23 Genera, 44 Species and 48 Strains in total. All these datasets from Bokulich et al. [26] are available at QIIME database² under their respective ID's. Since, these are simulated datasets the taxonomic profile of microbial organisms within them are known.

Real World Datasets:

4) *Canadian Soil*: The Canadian Soil dataset³ contains genomic data of soil spanning from Arctic Tundra to Agricultural soil suitable for different agricultural products.

5) *Body Sites*: This dataset contains composition of bacterial communities from up to 27 different body sites in healthy adults. A collection of 602 samples acquired from different body sites of human subjects are provided with meta-data.

6) *Global Soil*: The global soil data was taken from Ramirez et al. [29] which is a study of the below-ground diversity in New York City's Central Park.

B. Evaluation Metrics

We assess the performance of our developed clustering approach using the following commonly used metrics that are used for the assessment of (i) outputs from clustering algorithms, (ii) biodiversity within metagenomic samples and (iii) computational run time.

1) *Faiths phylogenetic diversity metric (PD)*: Faiths phylogenetic diversity [30] combines all the branch lengths of phylogenetic tree as a measure of diversity. So, if a new OTU is found and it is closely related to another OTU in the sample, it will contribute to a small increase to the PD score. However, if a new OTU from different lineage is found then it will contribute to a large increase in the PD score.

2) *Shannon Entropy*: Shannon-Wiener diversity index is defined as:

$$H = - \sum_{i=1}^s (p_i \log_2 p_i) \quad (1)$$

where s is the number of OTUs and p_i is the proportion of the community represented by OTU i . The Shannon index increases as both the richness and evenness of the community increase. The fact that the index incorporates both components of biodiversity can be seen as both a strength and a weakness. It is a strength because it provides a simple summary, but it is a weakness because it makes it difficult to compare communities that differ greatly in richness.

3) *Simpson's Index*: Simpsons index is defined as $1 - \text{dominance}$ or

$$1 - \sum p_i^2 \quad (2)$$

where p_i is the proportion of the community represented by OTU i . Simpsons index is based on the probability that any two individuals drawn at random from an infinitely large community belong to the same species. It measures *evenness* of the community from 0 to 1. Higher value of this index implies higher similarity and relatively lower diversity of microorganisms within a sample.

4) *F-measure*: In case of synthetic datasets, false-positive (FP; taxonomy/OTU string exists in observed but not expected), false-negative (FN; taxonomy/OTU exists in expected but not observed), and true-positive (TP; taxonomy/OTU exists in both observed and expected) measures were computed from cluster output and the ground truth

¹<https://github.com/ekopylova/otu-clustering>

²http://qiime.org/home_static/dataFiles.html

³<http://www.cm2bl.org/>

Table II
PERFORMANCE COMPARISON [F-SCORE AND PEARSON CORRELATION COEFFICIENT (ρ)]

Methods	Comparison Metric	Datasets					
		Synthetic			Real World		
		Bokulich ₂	Bokulich ₃	Bokulich ₆	Body Sites	Canadian Soil	Global Soil
UCLUST	F -Measure	0.39	0.40	0.51	N/A	N/A	N/A
CC _{UCLUST}	F -Measure	0.39	0.41	0.52	N/A	N/A	N/A
	(ρ) with Respect to UCLUST	0.9831	0.9753	0.9831	0.9682	0.8419	0.9824
SUMACLUSt	F -Measure	0.40	0.41	0.51	N/A	N/A	N/A
CC _{SUMACLUSt}	F -Measure	0.41	0.42	0.51	N/A	N/A	N/A
	(ρ) with Respect to SUMACLUSt	0.9709	0.9813	0.9538	0.9518	0.7643	0.8714
SWARM	F -Measure	0.46	0.48	0.55	N/A	N/A	N/A
CC _{SWARM}	F -Measure	0.46	0.49	0.56	N/A	N/A	N/A
	(ρ) with Respect to SWARM	0.9817	0.97861	0.9251	0.9648	0.7581	0.9143
CC _{MST}	F -Measure	0.41	0.42	0.52	N/A	N/A	N/A
	(ρ) with respect to UCLUST	0.7461	0.8143	0.8272	0.7417	0.6935	0.7974
	(ρ) with respect to SUMACLUSt	0.7182	0.6903	0.8613	0.8293	0.6791	0.8213
	(ρ) with respect to SWARM	0.7914	0.7352	0.8904	0.7619	0.7372	0.8502

Table shows values of F -measures and Pearson Correlation Coefficient (ρ -value) of UCLUST, SUMACLUSt, SWARM and their respective versions with Canopy clustering. F -measures is only available for synthetic datasets but not for real world datasets since no ground truths like known taxonomic profiles are available for them. ρ -value was calculated based on the taxonomy profiles at Genus level generated from clustered OTUs provided by a method and it's corresponding LSH based Canopy counterpart. Higher F -measures reflect better clustering by adhering to ground truth. Higher ρ -values reflect stronger correlation between taxonomic profiles.

which is the expected taxonomic composition. The following definitions were used:

$$precision = \frac{TP}{(TP + FP)} \quad (3)$$

$$recall = \frac{TP}{(TP + FN)} \quad (4)$$

$$FScore = \frac{2 \times precision \times recall}{(precision + recall)} \quad (5)$$

5) *Pearson Coefficient Correlation (ρ -value)*: After getting Operational Taxonomic Units (OTU) from a clustering method we create a taxonomic profile at the Genus level. Pearsons correlation coefficient was computed to measure the relatedness of taxonomic assignment between a pair of tools. Values range between -1 and 1, with -1 indicating a negative correlation, 0 indicating no correlation, and 1 indicating a positive correlation or strong relationship.

C. Parameter Settings

For kmers, the value of parameter K was set to 4. For Locality Sensitive Hashing the number of Hyperplanes (parameter d) was set to 150. For performance analysis and comparison we have chosen $T1 = 0.45$ and $T2 = 0.34$. We performed all the experiments on computers with Intel 5th generation Core i7 2.70GHz 64bit processor with 8 core CPUs and 12GB memory. For implementation we used Python 2.7.12 and QIIME [31] version 1.9.0 - a popular open source software that combines many metagenome clustering methods including the ones used in this work. Taxonomy for reported OTUs was assigned using the RDP Classifier [32] against the 97% representative databases for Greengenes [33] and Silva [34] for methods used in this study. We used

PyNast⁴ open source sequence aligner for aligning clustered output.

V. RESULTS DISCUSSION

A. Clustering Performance

Table II shows the performance of UCLUST [9], SUMACLUSt [11] and SWARM [10] and their corresponding versions with our proposed approach. Table II compares F -scores and Pearson Correlation Coefficient (ρ). F -scores are only available for synthetic benchmarks since taxonomic profile for them is known as ground truth. But F -scores are not available for real benchmarks as no ground truth is available for them. Correlation values were generated based on taxonomic profiles at Genus level generated from outputs from clustering methods. From Table II we can see that F -scores obtained from a clustering method and its corresponding Canopy clustering version are very similar and in some cases better. We see a higher F -score for Bokulich₂ benchmark from CC_{SUMACLUSt}. CC_{UCLUST} and CC_{SWARM} provided same F -scores as their respective naive versions for Bokulich₂. For Bokulich₃ benchmark we observed higher F -scores for all LSH based Canopy clustering methods. Finally for Bokulich₆ benchmark F -scores of UCLUST and SWARM were improved comparing to respective naive versions.

For all benchmarks we observe strong correlations between taxonomic profiles at genus level. The highest correlation was observed for Bokulich₆ benchmark between UCLUST and CC_{UCLUST} with 0.9831 and the lowest was observed for Canadian Soil metagenome data between SUMACLUSt and CC_{SUMACLUSt} with 0.7581. From

⁴<http://biocore.github.io/pynast/>

Table III
BIODIVERSITY COMPARISON [FAITHS PHYLOGENETIC DIVERSITY METRIC (PD), SHANNON AND SIMPSON]

Methods	Comparison Metric	Datasets					
		Synthetic			Real World		
		Bokulich ₂	Bokulich ₃	Bokulich ₆	Body Sites	Canadian Soil	Global Soil
UCLUST	PD Range	[171.95 – 221.85]	[186.90 – 212.84]	[104.51 – 104.51]	[1.46 – 46.79]	[0.30 – 1352.73]	[2.98 – 3.29]
	Shannon Range	[2.52 – 3.51]	[2.43 – 3.54]	[5.87 – 5.87]	[0.29 – 7.67]	[2.32 – 10.85]	[1.84 – 8.30]
	Simpson Range	[0.55 – 0.75]	[0.55 – 0.76]	[0.96 – 0.96]	[0.049 – 0.98]	[0.80 – 0.99]	[0.0 – 0.98]
CC _{UCLUST}	PD Range	[164.79 – 217.72]	[169.41 – 198.36]	[109.33 – 109.33]	[2.37 – 47.13]	[0.52 – 1419.31]	[3.02 – 3.81]
	Shannon Range	[2.61 – 3.83]	[2.92 – 3.91]	[6.41 – 6.41]	[0.93 – 7.13]	[3.26 – 7.61]	[3.72 – 7.38]
	Simpson Range	[0.64 – 0.87]	[0.56 – 0.93]	[0.96 – 0.96]	[0.081 – 0.99]	[0.84 – 0.99]	[0.21 – 0.99]
SUMACLUSt	PD Range	[106.00 – 162.78]	[142.85 – 174.19]	[89.22 – 89.22]	[0.93 – 39.47]	[0.59 – 1279.29]	[2.98 – 3.29]
	Shannon Range	[2.00 – 3.01]	[2.19 – 3.28]	[5.48 – 5.48]	[0.16 – 7.43]	[2.32 – 7.32]	[1.00 – 7.89]
	Simpson Range	[0.52 – 0.73]	[0.54 – 0.75]	[0.95 – 0.95]	[0.027 – 0.98]	[0.80 – 0.99]	[0.40 – 0.98]
CC _{SUMACLUSt}	PD Range	[114.96 – 171.57]	[147.85 – 187.91]	[93.81 – 93.81]	[0.86 – 41.63]	[0.81 – 1292.34]	[1.37 – 4.89]
	Shannon Range	[2.51 – 3.94]	[2.96 – 4.11]	[5.94 – 5.94]	[1.21 – 7.13]	[3.12 – 7.79]	[2.17 – 7.25]
	Simpson Range	[0.68 – 0.79]	[0.51 – 0.74]	[0.96 – 0.96]	[0.06 – 0.99]	[0.88 – 0.99]	[0.23 – 0.98]
SWARM	PD Range	[18.37 – 24.73]	[17.36 – 19.81]	[30.84 – 30.84]	[1.44 – 28.66]	[0.54 – 706.57]	[5.79 – 6.18]
	Shannon Range	[2.98 – 3.91]	[2.01 – 3.04]	[5.03 – 5.03]	[0.28 – 7.63]	[1.0 – 7.79]	[1.66 – 7.81]
	Simpson Range	[0.70 – 0.82]	[0.53 – 0.74]	[0.95 – 0.95]	[0.05 – 0.98]	[0.50 – 0.99]	[0.00 – 0.99]
CC _{SWARM}	PD Range	[19.18 – 26.87]	[18.43 – 22.61]	[31.48 – 31.48]	[2.34 – 29.97]	[1.37 – 748.71]	[2.34 – 8.46]
	Shannon Range	[1.66 – 4.87]	[1.19 – 4.13]	[6.03 – 6.03]	[0.89 – 7.13]	[2.81 – 8.06]	[2.81 – 7.87]
	Simpson Range	[0.66 – 0.88]	[0.41 – 0.81]	[0.91 – 0.91]	[0.11 – 0.99]	[0.74 – 0.99]	[0.14 – 0.99]
CC _{MST}	PD Range	[10.43 – 89.61]	[12.83 – 92.41]	[63.21 – 63.21]	[1.46 – 34.15]	[2.73 – 824.62]	[1.81 – 5.37]
	Shannon Range	[0.96 – 3.91]	[1.21 – 3.16]	[4.91 – 4.91]	[2.02 – 6.51]	[2.72 – 7.83]	[2.17 – 6.29]
	Simpson Range	[0.14 – 0.68]	[0.32 – 0.73]	[0.92 – 0.92]	[0.07 – 0.94]	[0.58 – 0.94]	[0.14 – 0.98]

Table shows ranges of values for Faiths Phylogeny Diversity (PD), Shannon and Simpson coefficient over all samples in a dataset in the format $[minimum - maximum]$. Most of these datasets contain multiple samples and Alpha diversity metrics like PD, Shannon and Simpson values are generated for each of these samples separately. Biodiversity metric values changes significantly over samples e.g. diversity from hair samples and teeth cavity are supposedly different. So instead of mean values this Table represents $[minimum - maximum]$ ranges of values a sample can take. Similar ranges reflect similar diversity.

Table IV
RUNTIME COMPARISON (IN MINUTES)

Datasets		Methods											
Type	Title # of Reads	UCLUST	CC _{UCLUST}	Speed Up	SUMACLUSt	CC _{SUMACLUSt}	Speed Up	SWARM	CC _{SWARM}	Speed Up	MST	CC _{MST}	Speed Up
Synthetic	Bokulich ₂ 6,938,836	12.71	6.03	2.10x	87.53	33.03	2.65x	128.12	37.87	3.38x	564.60	294.36	1.91x
	Bokulich ₃ 3,594,237	8.91	4.89	1.82x	11.73	7.89	1.48x	9.10	6.93	1.31x	369.7	204.61	1.80x
	Bokulich ₆ 250,903	1.21	1.08	1.12x	1.28	1.34	0.96	1.97	1.29	1.53x	252.81	146.41	1.73x
Real World	Body Sites 886,630	2.12	1.51	1.40x	15.42	7.88	1.95x	3.64	1.71	2.13x	307.2	174.27	1.76x
	Canadian Soil 2,966,053	9.55	5.91	1.62x	363.96	81.18	4.48x	97.50	32.16	3.03x	472.43	254.91	1.85x
	Global Soil 9,252,764	72.47	43.53	1.67x	510.92	66.27	7.70x	269.07	56.14	4.79x	988.93	464.29	2.13x

observed F-scores and correlation values we can say that running a sequence clustering method with our proposed LSH based Canopy cluster framework can bring similar or better results. CC_{MST} provided relatively similar F-scores and high ρ values indicating positive correlations of taxonomic outcomes with other methods.

B. Biodiversity Comparison

Clustering sequences in metagenome data will provide OTUs that represent biodiversity contained in the samples from which data is collected. Hence we also need to compare the biodiversity represented by OTUs from naive methods and their respective Canopy clustering variations used in this study. Table III shows Faiths phylogenetic diversity metric (PD), Shannon and Simpson index after clustering with different methods. These metrics are some of the popular Alpha Diversity metrics meaning that they measure species diversity in sites or habitats at a local scale. These metric values are generated per sample basis. Table III shows ranges of metric values in $[minimum - maximum]$ format

provided by methods with and without Canopy clustering. Any sample of a dataset will take value from this range. These ranges may not be same since OTUs vary over clustering methods. But they should be similar for a method and it's Canopy counterpart. We can see from Table III that Canopy clustering based methods produce similar ranges of values as their naive counterparts. No significant changes in diversity metric values were observed which indicates that Canopy based approaches reproduce similar biodiversity while reducing runtime.

C. Runtime Comparison

Table IV shows the runtime in minutes of UCLUST, SUMACLUSt, SWARM, MST and their respective versions with our proposed Canopy clustering pipeline. Improving runtime while maintaining similar or better clustering results was one of the major motivations of our proposed approach. We can see from Table IV that CC_{UCLUST} outperforms UCLUST for all datasets. The highest speed up for UCLUST was observed for Bokulich₂ dataset where CC_{UCLUST} was




Figure 2. Effect of number of multiprocess on Runtime of CC_{UCLUST} , $CC_{SUMACLUST}$, CC_{SWARM} and CC_{MST} for Bokulich₂ and GLocal Soil datasets, two of the largest datasets used in this study.

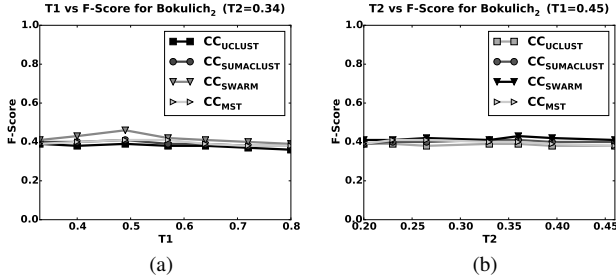


Figure 3. 3a and 3b show effect of varying $T1$ and $T2$ on F -scores for the largest synthetic dataset Bokulich₂

2.10 times faster than UCLUST. $CC_{SUMACLUST}$ outperforms naive SUMACLUST in most cases. Major improvement in runtime was observed in large scale real world metagenome datasets specially Canadian Soil and Global Soil datasets. The highest speed-up for SUMACLUST with Canopy clustering was observed to be 7.70x for Global Soil dataset. We can see a similar footprints for CC_{SWARM} and CC_{MST} from Table IV. The highest speedup in runtime for SWARM and MST with Canopy clustering was observed to be 4.79x and 2.13x respectively for Global Soil dataset which is the largest dataset used in this study. From these observations we can conclude that LSH based Canopy clustering not only reduces the runtime of an expensive clustering method but also scales well and performs better for larger benchmark datasets comparing to the smaller ones.

D. Effect of Varying Number of Multi-processes

Figure 2a-2b shows how the number of multiprocess can affect runtime of CC_{UCLUST} , $CC_{SUMACLUST}$, CC_{SWARM} and CC_{MST} on two largest of the benchmarks used in this study titled Bokulich₂ and Global Soil. We observe that increasing the number of multiprocess can reduce total runtime. Steeper curves show reduction in runtime with increasing number of multiprocess. Significant reductions in run time were observed CC_{MST} , $CC_{SUMACLUST}$ and $CC_{SUMACLUST}$ compared to CC_{UCLUST} .

E. Effect of Varying $T1$ and $T2$ Parameter of Canopy Clustering

Figure 3a-3b shows effect of varying $T1$ ($T2$ fixed at 0.34) and $T2$ ($T1$ fixed at 0.45) on F -scores for the largest synthetic dataset Bokulich₂. Reducing $T1$ leads to comparatively *strict* soft-threshold which will reduce repetitions of instances in multiple canopies. For a fixed $T2$ this implies that lower $T1$ will yield better canopy assignment. From Figure 3a we can say that when $T1$'s range is in 0.4 to 0.6 our proposed approach provides better F -Scores. On the other hand, increasing $T2$ leads to comparatively *relaxed* tight-threshold for canopies. As a result instances will be prematurely assigned to canopies without waiting for best match. From Figure 3b we can say that when $T2$'s range is in 0.2 to 0.35 our proposed approach achieves better F -Scores. Lower $T2$ may lead to higher runtime since instances will continue to reappear until the $T2$ is met.

VI. CONCLUSION AND FUTURE WORK

We propose a framework that can be used as pre-clustering for any accurate and relatively expensive clustering on large scale metagenomic datasets. Our approach takes advantage of the multi-core CPU systems by partitioning the large dataset roughly with fast and cheaper pairwise distance measure and then deploying comparatively expensive clustering in parallel which considers only data points that are inside a partition. Our proposed approach scales well with large datasets and provide significant reduction in computation time. We demonstrate that our approach provides similar outcome in terms of biodiversity metrics, ground truth and taxonomic correlation with corresponding expensive clustering methods. Standalone metagenomic clustering algorithm with canopy requires future investigation in future.

REFERENCES

- [1] Qin et al., "A human gut microbial gene catalogue established by metagenomic sequencing," *nature*, vol. 464, no. 7285, pp. 59–65, 2010.
- [2] Turnbaugh et al., "The human microbiome project: exploring the microbial part of ourselves in a changing world," *Nature*, vol. 449, no. 7164, p. 804, 2007.
- [3] Myers et al., "A whole-genome assembly of drosophila," *Science*, vol. 287, no. 5461, pp. 2196–2204, 2000.
- [4] A. Charuvaka and H. Rangwala, "Evaluation of short read metagenomic assembly," *BMC genomics*, vol. 12, no. 2, p. 1, 2011.
- [5] Chakravorty et al., "A detailed analysis of 16s ribosomal rna gene segments for the diagnosis of pathogenic bacteria," *Journal of microbiological methods*, vol. 69, no. 2, pp. 330–339, 2007.
- [6] Blaxter et al., "Defining operational taxonomic units using dna barcode data," *Philosophical Transactions of the Royal Society B: Biological Sciences*, vol. 360, no. 1462, pp. 1935–1943, 2005.

- [7] T. S. Schmidt, J. F. M. Rodrigues, and C. Von Mering, "Ecological consistency of ssu rna-based operational taxonomic units at a global scale," *PLoS Comput Biol*, vol. 10, no. 4, p. e1003594, 2014.
- [8] Kopylova et al., "Open-source sequence clustering methods improve the state of the art," *mSystems*, vol. 1, no. 1, 2016.
- [9] Edgar, Robert C., "Search and clustering orders of magnitude faster than blast," *Bioinformatics*, vol. 26, no. 19, pp. 2460–2461, 2010.
- [10] M. et al., "Swarm v2: highly-scalable and high-resolution amplicon clustering," *PeerJ*, vol. 3, p. e1420, 2015.
- [11] C, Mercier and F, Boyer and A, Bonin and E, Coissac, "Sumatra and sumacust: fast and exact comparison and clustering of sequences [submitted for publication]," Available at <https://git.metabarcoding.org/obitools/sumatra/wikis/home>, 2014.
- [12] S. et al., "Introducing mothur: Open-source, platform-independent, community-supported software for describing and comparing microbial communities," *Applied and Environmental Microbiology*, vol. 75, no. 23, pp. 7537–7541, 2009.
- [13] Li, Weizhong and Godzik, Adam, "Cd-hit: a fast program for clustering and comparing large sets of protein or nucleotide sequences," *Bioinformatics*, vol. 22, no. 13, pp. 1658–1659, 2006.
- [14] Z. Rasheed, H. Rangwala, and D. Barbara, "Efficient clustering of metagenomic sequences using locality sensitive hashing," in *SDM*. SIAM, 2012, pp. 1023–1034.
- [15] Z. Rasheed and H. Rangwala, *MC-MinH: Metagenome Clustering using Minwise based Hashing*, ch. 74, pp. 677–685.
- [16] A. Z. Broder, M. Charikar, A. M. Frieze, and M. Mitzenmacher, "Min-wise independent permutations," *Journal of Computer and System Sciences*, vol. 60, no. 3, pp. 630–659, 2000.
- [17] O. et al., "Mash: fast genome and metagenome distance estimation using minhash," *bioRxiv*, p. 029827, 2016.
- [18] O. Tanaseichuk, J. Borneman, and T. Jiang, "Separating metagenomic short reads into genomes via clustering," *Algorithms for Molecular Biology*, vol. 7, no. 1, p. 1, 2012.
- [19] Y.-W. Wu and Y. Ye, "A novel abundance-based algorithm for binning metagenomic sequences using 1-tuples," *Journal of Computational Biology*, vol. 18, no. 3, pp. 523–534, 2011.
- [20] S. Chatterji, I. Yamazaki, Z. Bai, and J. A. Eisen, "Compostbin: A dna composition-based algorithm for binning environmental shotgun reads," in *Annual International Conference on Research in Computational Molecular Biology*. Springer, 2008, pp. 17–28.
- [21] K. et al., "Open-source sequence clustering methods improve the state of the art," *mSystems*, vol. 1, no. 1, pp. e00003–15, 2016.
- [22] McCallum, Andrew and Nigam, Kamal and Ungar, Lyle H., "Efficient clustering of high-dimensional data sets with application to reference matching," in *Proceedings of the Sixth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, ser. KDD '00. New York, NY, USA: ACM, 2000, pp. 169–178.
- [23] Gionis, Aristides and Indyk, Piotr and Motwani, Rajeev, "Similarity search in high dimensions via hashing," in *Proceedings of the 25th International Conference on Very Large Data Bases*, ser. VLDB '99. San Francisco, CA, USA: Morgan Kaufmann Publishers Inc., 1999, pp. 518–529.
- [24] P. K. Jana and A. Naik, "An efficient minimum spanning tree based clustering algorithm," pp. 1–5, 2009.
- [25] M. Laszlo and S. Mukherjee, "Minimum spanning tree partitioning algorithm for microaggregation," *IEEE Transactions on Knowledge and Data Engineering*, vol. 17, no. 7, pp. 902–911, 2005.
- [26] Bokulich, NA and Subramanian, S and Faith JJ and Gevers, D and Gordon, JI and Knight, R and Mills, DA and Caporaso, JG, "Quality-filtering vastly improves diversity estimates from illumina amplicon sequencing," *Nat Methods*, vol. 10, pp. 57–59, 2013.
- [27] Neufeld et al., "Open resource metagenomics: a model for sharing metagenomic libraries," *Standards in Genomic Sciences*, vol. 5, pp. 203–210, 2011.
- [28] Costello, Elizabeth K. and Lauber, Christian L. and Hamady, Micah and Fierer, Noah and Gordon, Jeffrey I. and Knight, Rob, "Bacterial community variation in human body habitats across space and time," *Science*, vol. 326, no. 5960, pp. 1694–1697, 2009.
- [29] Ramirez et al., "Biogeographic patterns in below-ground diversity in new york city's central park are similar to those observed globally," *Proceedings of the Royal Society of London B: Biological Sciences*, vol. 281, no. 1795, 2014.
- [30] Faith, Daniel P, "Conservation evaluation and phylogenetic diversity," *Biological conservation*, vol. 61, no. 1, pp. 1–10, 1992.
- [31] Caporaso et al., "Qiime allows analysis of high-throughput community sequencing data," *Nat Meth*, vol. 7, no. 5, pp. 335–336, 2010.
- [32] Q. Wang, G. M. Garrity, J. M. Tiedje, and J. R. Cole, "Naive bayesian classifier for rapid assignment of rna sequences into the new bacterial taxonomy," *Applied and environmental microbiology*, vol. 73, no. 16, pp. 5261–5267, 2007.
- [33] DeSantis et al., "Greengenes, a chimera-checked 16s rna gene database and workbench compatible with arb," *Applied and environmental microbiology*, vol. 72, no. 7, pp. 5069–5072, 2006.
- [34] Pruesse et al., "Silva: a comprehensive online resource for quality checked and aligned ribosomal rna sequence data compatible with arb," *Nucleic acids research*, vol. 35, no. 21, pp. 7188–7196, 2007.